

Storing and Managing Data with Redis and Apache Kafka on Heroku

PROVISION APACHE KAFKA ON HEROKU



Saravanan Dhandapani

SOFTWARE ARCHITECT

@dsharu



Course Structure

Part 1

Apache Kafka on Heroku

Part2


Redis on Heroku




Provision Apache Kafka on Heroku



Globomantics E-Commerce Portal



LinkedIn Issue



LinkedIn Issue

1 trillion messages per day

Peak volume of 10 million messages per second

More than 400 millions users

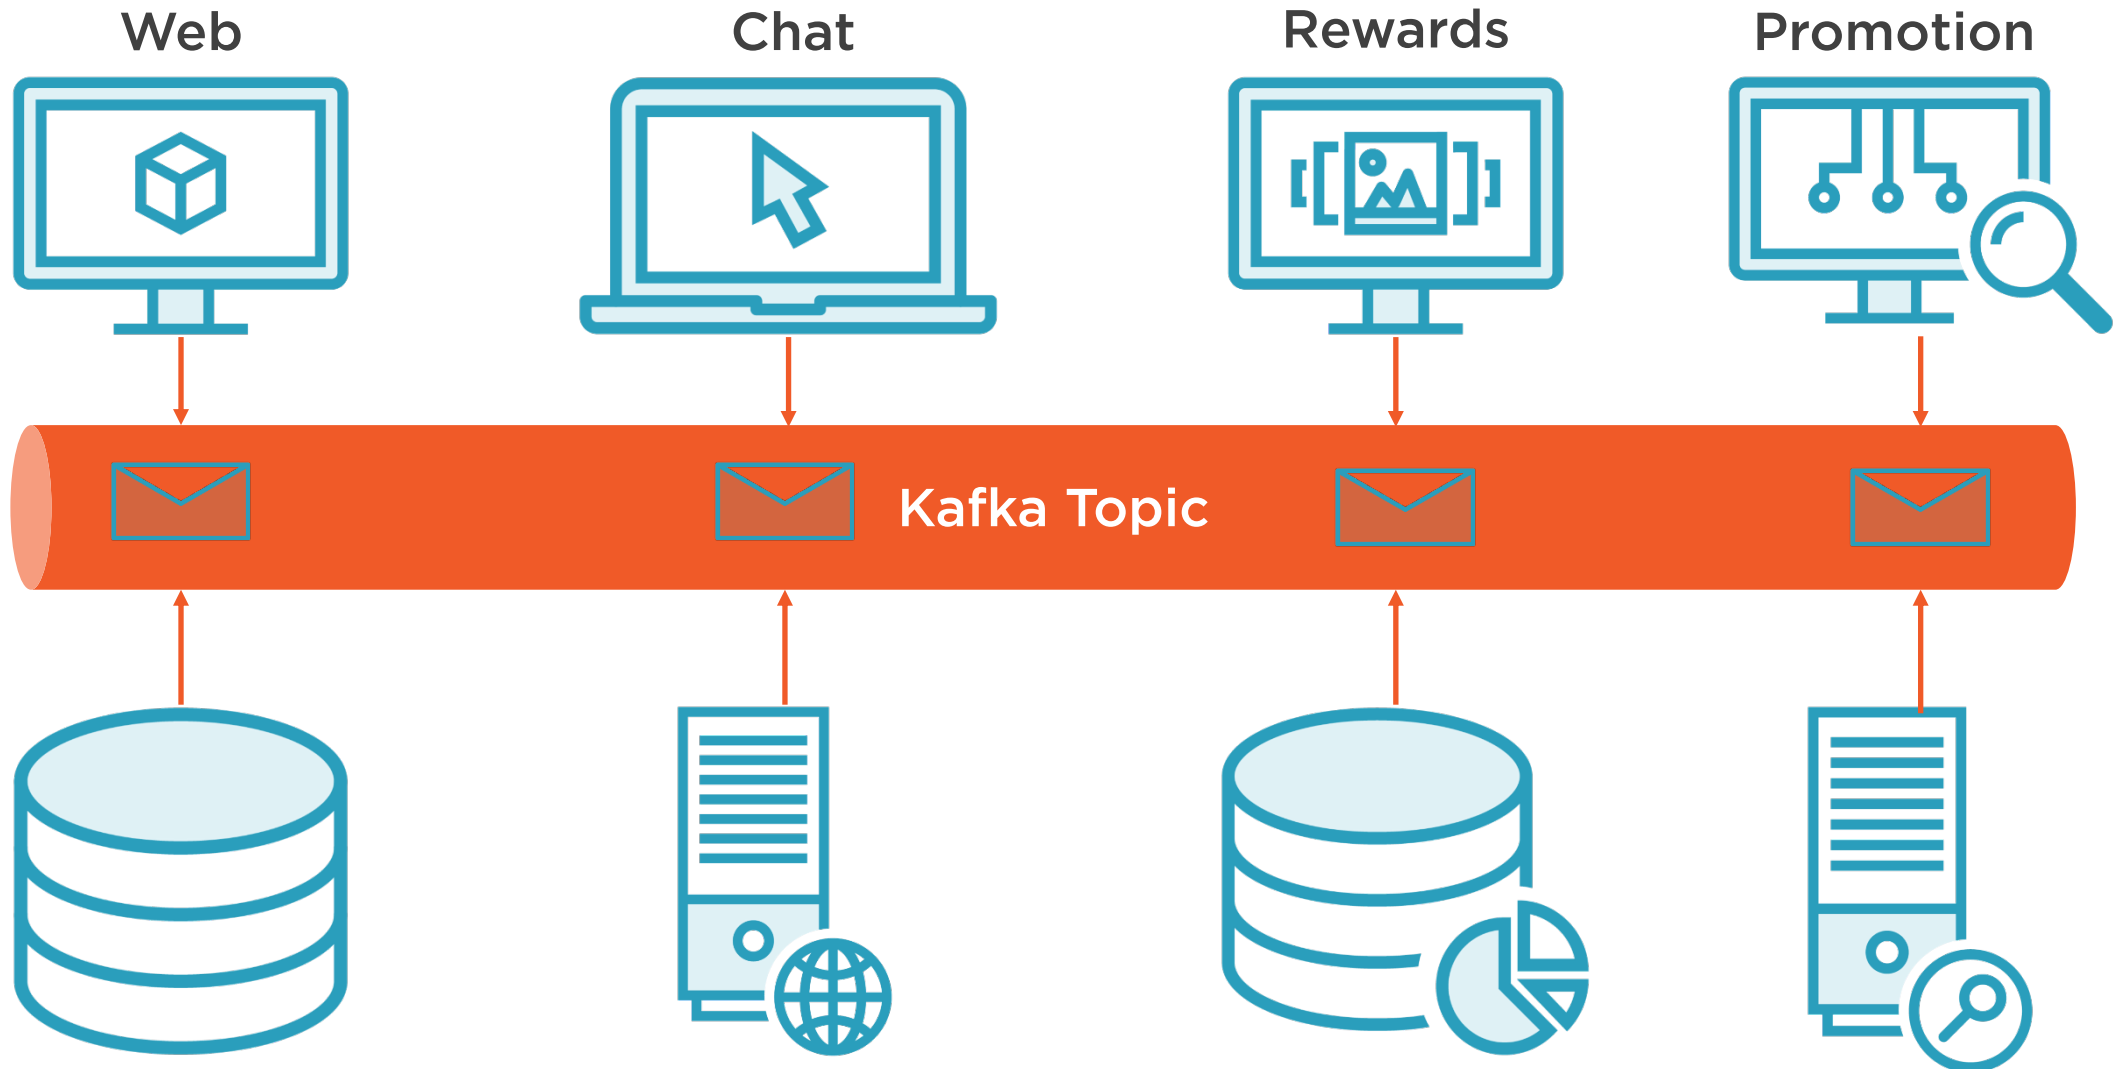
Independently created pipelines routing this data

Kafka



- Open sourced in 2012



Globomantics E-Commerce Portal



Postal Analogy



Core Kafka Concepts

Kafka brokers

- Servers running Kafka topics

Kafka cluster

Zookeeper

- Manages the brokers in the cluster

Producers and Consumers

Topics

- Message containers

Partitions

Configuration options

- Message retention
- Replication factor



Why Kafka? Why Heroku?



Kafka Business Cases

Smart devices

Process
streaming data

**Streaming
applications**

Infrastructure
and application
errors

Publish metrics
and create
dashboards

**Monitoring and
Alerting**

Replay events
and determine
root cause

**Complex error
handling**

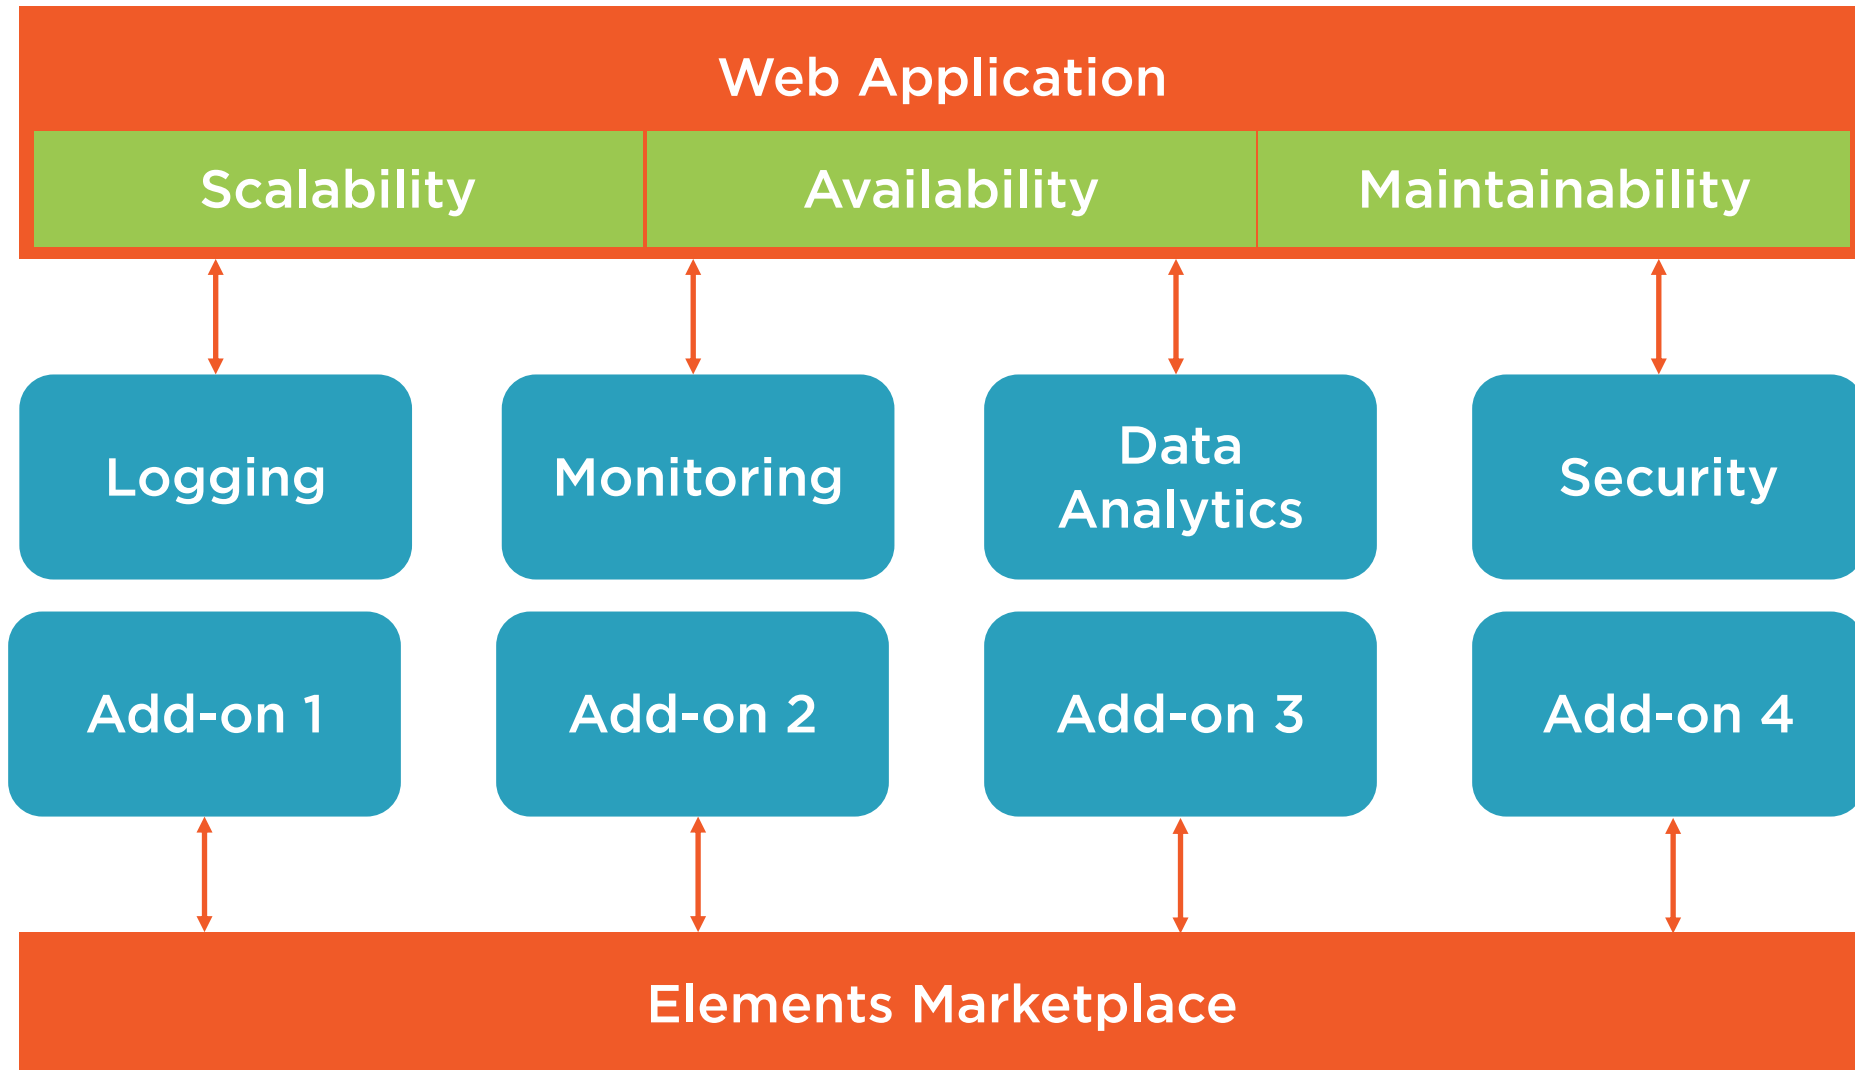
Microservice
architecture

Service to
service
communication

**Promote loose
coupling**



Heroku Architecture



Why Heroku?



Scale based on your needs



Resources are highly available and addressable



Maintain external configuration variables



Rich CLI experience to provision and monitor



Stream PHI and PII data



Provision Apache Kafka Addon



Kafka Basic Plan Comparison

| Plan name | Producer quota (KB/s) | Consumer quota (KB/s) | Capacity (GB) | Max partitions | Price/month |
|-----------|-----------------------|-----------------------|---------------|----------------|-------------|
| Basic-0 | 64 | 256 | 4 | 240 | \$100 |
| Basic-1 | 512 | 2048 | 32 | 480 | \$125 |
| Basic-2 | 4096 | 16384 | 64 | 960 | \$175 |



Kafka Basic Plan Use Cases

Limited production usage

Development and testing environments

Experimentation and prototyping purposes



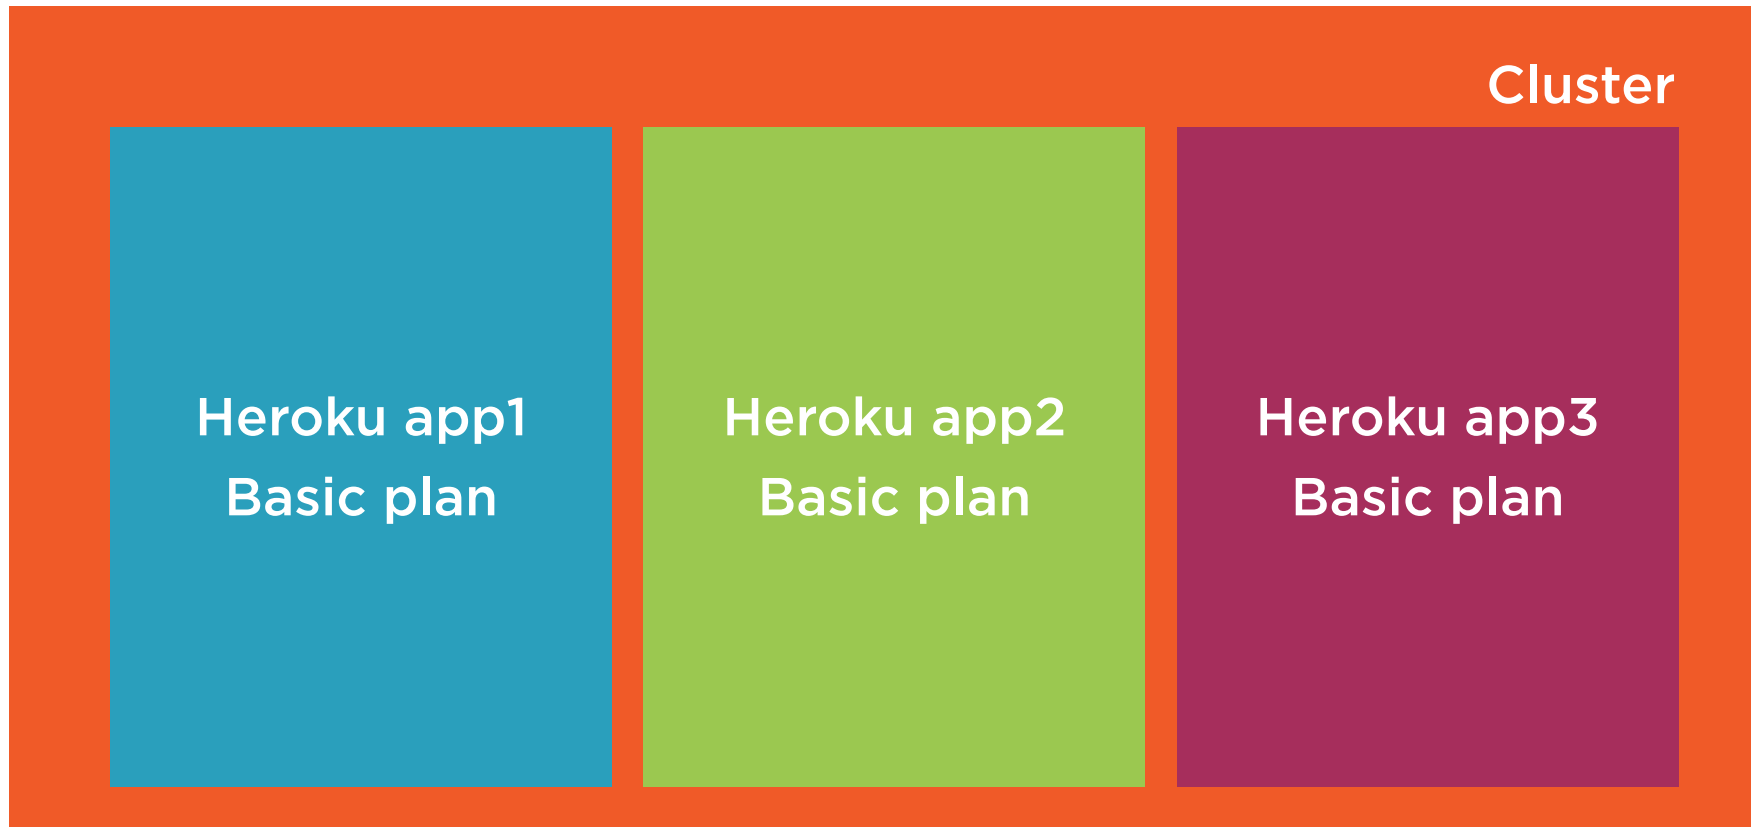
Multi-Tenant Kafka on Heroku



Multi-Tenancy



Heroku Basic Plan



Dedicated Cluster

A minimum of 3 Kafka brokers

2 weeks of data retention

Price starting from \$1500/month

Unique prefix not needed



Features
unavailable in
Basic Plan

Staging environment testing

- heroku kafka:fail

Upgrading environment

- heroku kafka:upgrade



Connect and Manage Kafka on Heroku

